

The data

Today we will work on [2018 Data Science Bowl](#) dataset.

You can download images and masks directly from the url or using Kagge API :

```
kaggle competitions download -c data-science-bowl-2018
```

After downloading the data, unpack them and move to preferred destination. For this example we will be interested only in `stage1_train` and `stage1_test` subdirectories, so you can delete other files if you want.

Before we start, let's investigate a little bit.

```
library(tidyverse)
library(platypus)
library(abind)
library(here)

# Print current working directory
here()

# [1] "/home/majul16/Desktop/PROJECTS/Moje Projekty/platypus"

# Set directories with the data and models
data_path <- here("examples/data/data-science-bowl-2018/")
models_path <- here("examples/models/")

# Investigate one instance of data (image + masks)
sample_image_path <- here("examples/data/data-science-bowl-2018/stage1_train/00071198d059ba7f5914a526d124d28e6d010c92466da21d4a04cd5413362552/")

list.files(sample_image_path, full.names = TRUE) %>%
  set_names(basename(.)) %>%
  map(~ list.files(.))

# $images
# [1] "00071198d059ba7f5914a526d124d28e6d010c92466da21d4a04cd5413362552.png"
#
# $masks
# [1] "07a9bf1d7594af2763c86e93f05d22c4d5181353c6d3ab30a345b908ffe5aad5.png"
# [2] "0e548d0af63ab451616f082eb56bde13eb71f73dfda92a03f8e8ad42ebb4881.png"
# [3] "0ea1f9e30124e4aef1407af239ff42fd6f5753c09b4c5cac5d08023c328d7f05.png"
# [4] "0f5a3252d05ecdf453bdd5e6ad5322c454d8ec2d13ef0f0bf45a6f6db45b5639.png"
# [5] "2c47735510ef91a11fde42b317829cee5fc04d05a797b90008803d7151951d58.png"
# [6] "4afa39f2a05f9884a5ff030d678c6142379f99a5baaf4f1ba7835a639cb5"
```

```
0751.png"
# [7] "4bc58dbdefb2777392361d8b2d686b1cc14ca310e009b79763af46e853e6
c6ac.png"
# [8] "4e3b49fb14877b63704881a923365b68c1def111c58f23c66daa49fef4b6
32bf.png"
# [9] "5522143fa8723b66b1e0b25331047e6ae6eeec664f7c8abeba687e0de0f9
060a.png"
# [10] "58656859fb9c13741eda9bc753c3415b78d1135ee852a194944dee88ab70
acf4.png"
# [11] "6442251746caac8fc255e6a22b41282ffcfafebadbd240ee0b604808ff9e
3383.png"
# [12] "7ff04129f8b6d9aaf47e062eadce8b3fcff8b4a29ec5ad92bca926ac2b72
63d2.png"
# [13] "8bbec3052bcec900455e8c7728d03facb46c880334bcc4fb0d1d066dd6c7
c5d2.png"
# [14] "9576fe25f4a510f12eecbabfa2e0237b98d8c2622b9e13b9a960e2afe6da
844e.png"
# [15] "95deddb72b845b1a1f81a282c86e666045da98344eaa2763d67e2ab80bc2
e5c3.png"
# [16] "a1b0cdb21f341af17d86f23596df4f02a6b9c4e0d59a7f74aaf28b9e408a
4e4c.png"
# [17] "aa154c70e0d82669e9e492309bd00536d2b0f6eeec1210014bbafbfc554b
377c.png"
# [18] "acba6646e8250aab8865cd652dfaa7c56f643267ea2e774aee97dc2342d8
79d6.png"
# [19] "ae00049dc36a1e5ffaafcdeadb44b18a9cd6dfd459ee302ab041337529bd4
1cf2.png"
# [20] "af4d6ff17fa7b41de146402e12b3bab1f1fe3c1e6f37da81a54e002168b1
e7dd.png"
# [21] "b0cbc2c553f9c4ac2191395236f776143fb3a28fb77b81d3d258a2f45361
ca89.png"
# [22] "b6fc3b5403de8f393ca368553566eaf03d5c07148539bc6141a486f1d185
f677.png"
# [23] "be98de8a7ba7d5d733b1212ae957f37b5b69d0bf350b9a5a25ba4346c29e
49f7.png"
# [24] "cb53899ef711bce04b209829c61958abdb50aa759f3f896eb7ed868021c2
2fb4.png"
# [25] "d5024b272cb39f9ef2753e2f31344f42dd17c0e2311c4927946bc5008d29
5d2e.png"
# [26] "f6eee5c69f54807923de1ceb1097fc3aa902a6b20d846f111e806988a426
9ed0.png"
# [27] "ffae764df84788e8047c0942f55676c9663209f65da943814c6b3aca78d8
e7f7.png"
```

As you can see each image has its own directory, that has two subdirectories inside:

- **images** – contains original image that will be the input of the neural network
- **masks** – contains **one or more** segmentation masks. **Segmentation mask** is simply telling us which pixel belongs to which class, and this is what we will try to predict.

For the modeling, beside **train** and **test** sets, we will also need a **validation** set (No one is forcing you, but it's a good practice!):

```

train_path <- here("examples/data/data-science-bowl-2018/stage1_train/")
test_path <- here("examples/data/data-science-bowl-2018/stage1_test/")
validation_path <- here("examples/data/data-science-bowl-2018/stage1_validation/")

if (!dir.exists(validation_path)) {
  dir.create(validation_path)
  # List train images
  train_samples <- list.files(train_path, full.names = TRUE)
  set.seed(1234)
  # Select 10% for validation
  validation_samples <- sample(train_samples, round(0.1 *
length(train_samples)))
  validation_samples %>%
    walk(~ system(paste0('mv "', ., '" "', validation_path, '"')))
}

```

Semantic segmentation with U-Net

Since we now something about our data, we can now move to the modeling part. We will start by selecting the architecture of the neural network. In case of semantic segmentation there is a few different choices like **U-Net**, **Fast-FCN**, **DeepLab** and many more. For the time being in the [platypus](#) package you have access only to the **U-Net** architecture.




Fig. 1. U-net architecture (example for 32x32 pixels in the lowest resolution). Each blue box corresponds to a multi-channel feature map. The number of channels is denoted on top of the box. The x-y-size is provided at the lower left edge of the box. White boxes represent copied feature maps. The arrows denote the different operations.

U-Net was originally developed for biomedical data segmentation. As you can see in the picture above architecture is very similar to autoencoder and it looks like the letter **U**, hence the name. Model is composed of 2 parts, and each part has some number of **convolutional blocks** (3 in the image above). Number of blocks will be hyperparameter in our model.

To build a **U-Net** model in [platypus](#) use `u_net` function. You have to specify:

- number of convolutional blocks,
- input image height and width – must be in the form 2^N !,
- will input image be loaded as grayscale or RGB,
- number of classes – in our case we have only 2 (background and nuclei)
- additional arguments form CNN like number of filters, dropout rate

```
blocks <- 4 # Number of U-Net convolutional blocks
n_class <- 2 # Number of classes
net_h <- 256 # Must be in a form of 2^N
net_w <- 256 # Must be in a form of 2^N
grayscale <- FALSE # Will input image be in grayscale or RGB

DCB2018_u_net <- u_net(
  net_h = net_h,
  net_w = net_w,
  grayscale = grayscale,
  blocks = blocks,
  n_class = n_class,
  filters = 16,
  dropout = 0.1,
  batch_normalization = TRUE,
  kernel_initializer = "he_normal"
)
```

After that it's time to select **loss** and additional metrics. Because semantic segmentation is in essence classification for each pixel instead of the whole image, you can use **categorical cross-entropy** as a loss function and **accuracy** as a metric. Other common choice, available in `platypus`, would be **dice coefficient/loss**. You can think of it as of a **F1-metric** for semantic segmentation.

```
DCB2018_u_net %>%
  compile(
    optimizer = optimizer_adam(lr = 1e-3),
    loss = loss_dice(),
    metrics = metric_dice_coeff()
  )
```

The next step will be data ingestion. As you remember we have a separate directory and multiple masks for each image. That's not a problem for `platypus`! You can ingest data using `segmentation_generator` function. The first argument to specify is the directory with all the images and masks. To tell `platypus` that it has to load images and masks from separate directories for each data sample specify argument `mode = "nested_dirs"`. Additionally you can set images/masks subdirectories names using `subdirs` argument. `platypus` will automatically merge multiple masks for each image, but we have to tell him how to recognize which pixel belongs to which class. In the segmentation masks each class is recognized by a specific RGB value. In our case we have only black (R = 0, G = 0, B = 0) pixel for background and white (R = 255, G = 255, B = 255) pixels for nuclei. To tell `platypus` how to recognize classes on segmentation masks use `colormap` argument.

```
binary_colormap

# [[1]]
```

```

# [1] 0 0 0
#
# [[2]]
# [1] 255 255 255

train_DCB2018_generator <- segmentation_generator(
  path = train_path, # directory with images and masks
  mode = "nested_dirs", # Each image with masks in separate folder
  colormap = binary_colormap,
  only_images = FALSE,
  net_h = net_h,
  net_w = net_w,
  grayscale = FALSE,
  scale = 1 / 255,
  batch_size = 32,
  shuffle = TRUE,
  subdirs = c("images", "masks") # Names of subdirs with images and masks
)

# 603 images with corresponding masks detected!
# Set 'steps_per_epoch' to: 19

validation_DCB2018_generator <- segmentation_generator(
  path = validation_path, # directory with images and masks
  mode = "nested_dirs", # Each image with masks in separate folder
  colormap = binary_colormap,
  only_images = FALSE,
  net_h = net_h,
  net_w = net_w,
  grayscale = FALSE,
  scale = 1 / 255,
  batch_size = 32,
  shuffle = TRUE,
  subdirs = c("images", "masks") # Names of subdirs with images and masks
)

# 67 images with corresponding masks detected!
# Set 'steps_per_epoch' to: 3

```

We can now fit the model.

```

history <- DCB2018_u_net %>%
  fit_generator(
    train_DCB2018_generator,
    epochs = 20,
    steps_per_epoch = 19,
    validation_data = validation_DCB2018_generator,
    validation_steps = 3,
    callbacks = list(callback_model_checkpoint(
      filepath = file.path(models_path, "DSB2018_w.hdf5"),
      save_best_only = TRUE,
      save_weights_only = TRUE,

```

```

        monitor = "dice_coeff",
        mode = "max",
        verbose = 1)
    )
)

```

And calculate predictions for the new images. Our model will return a 4-dimensional array (number of images, height, width, number of classes). Each pixel will have N probabilities, where N is number of classes. To transform raw predictions into segmentation map (by selecting class with max probability for each pixel) you can use `get_masks` function.

```

test_DCB2018_generator <- segmentation_generator(
  path = test_path,
  mode = "nested_dirs",
  colormap = binary_colormap,
  only_images = TRUE,
  net_h = net_h,
  net_w = net_w,
  grayscale = FALSE,
  scale = 1 / 255,
  batch_size = 32,
  shuffle = FALSE,
  subdirs = c("images", "masks")
)

# 65 images detected!
# Set 'steps_per_epoch' to: 3

test_preds <- predict_generator(DCB2018_u_net, test_DCB2018_generator, 3)
dim(test_preds)

# [1] 65 256 256 2

test_masks <- get_masks(test_preds, binary_colormap)
dim(test_masks[[1]])

# [1] 256 256 3

```

To visualize predicted masks with the original images you can use `plot_masks` function.

```

test_imgs_paths <- create_images_masks_paths(test_path, "nested_dirs",
FALSE, c("images", "masks"), ";")$images_paths

plot_masks(
  images_paths = test_imgs_paths[1:4],
  masks = test_masks[1:4],
  labels = c("background", "nuclei"),
  colormap = binary_colormap
)

```

